

DAILY ASSESSMENT FORMAT

Date:	17 JUNE 2020	Name:	PAVITHRAN S
Course:	STATISTICAL ANALYSIS	USN:	4AL17EC068
Topic:	STATISTICAL ANALYSIS	Semester & Section:	6 TH B
Github Repository:	Pavithran		

FORENOON SESSION DETAILS

Image of session

The screenshot shows a web interface for a Great Learning course. At the top, there's a navigation bar with 'Home', 'Live Sessions', and 'Certificates'. A 'My Courses' button and a user profile icon are on the right. Below the navigation bar, the breadcrumb trail reads 'Courses / Statistical Learning / Normal Distribution'. The main content area is divided into a left sidebar and a right main panel. The sidebar has a 'Content' section with a list of topics: 'Agenda', 'Case study on statistics and Probability Theory', 'Solution for case study', 'Introduction to Probability', 'Rules for Probability calculation', 'Bayes' Theorem', and 'Normal Distribution' (which is currently selected). Below this are sections for 'Learning Material', 'Quiz', and 'Claim Your Course Certificate'. The main panel displays the title 'Normal Distribution' and a video player. The video shows a normal distribution curve on a green background with a red line. A label 'Mean Median Mode' points to the peak of the curve. Below the video, the name 'DR. P.K. VISWANATHAN' is displayed. The video player controls at the bottom show a progress bar at 0:12, a volume icon, and a full screen icon.

Report – Report can be typed or hand written for up to two pages.

What is Statistical Learning?

Let, suppose that we observe a response Y and p different predictors $X = (X_1, X_2, \dots, X_p)$. In general, we can say:

$$Y = f(X) + \varepsilon$$

Here f is an unknown function, and ε is the *random error term*.

In essence, statistical learning refers to a set of approaches for estimating f .

In cases where we have set of X readily available, but the output Y , not so much, the error averages to zero, and we can say:

$$\hat{Y} = f(X)$$

where \hat{f} represents our estimate of f and \hat{Y} represents the resulting prediction.

Hence for a set of predictors X , we can say:

$$E(Y - \hat{Y})^2 = E[f(X) + \varepsilon - f(X)]^2 \Rightarrow E(Y - \hat{Y})^2 = [f(X) - \hat{f}(X)]^2 + \text{Var}(\varepsilon)$$

where,

- $E(Y - \hat{Y})^2$ represents *the expected value* of the squared difference between actual and expected result.
- $[f(X) - \hat{f}(X)]^2$ represents the **reducible error**. It is reducible because we can potentially improve the accuracy of \hat{f} by better modeling.
- $\text{Var}(\varepsilon)$ represents the **irreducible error**. It is irreducible because no matter how well we estimate f , we cannot reduce the error introduced by *variance* in ε .

Regression Vs Classification Problem

Variables, Y , can be broadly be characterized as *quantitative* or *qualitative* (also known as *categorical*). Quantitative variables take on numerical values, e.g., age, height, income, price, and much more. Estimating qualitative responses is often termed as a **regression problem**. Qualitative variables take on categorical values, e.g., gender, brand, parts of speech, and much more. Estimating qualitative responses is often termed as a **classification problem**.

There is no free lunch in statistics: no one method dominates all other over all possible data sets.

Variance And Bias

Variance refers to the amount by which f would change if we estimated with different training data sets. In general, when we over-fit a model on a given training data set (reducible error in training set is very low but on test set is very high), we get a model that has higher variance since any change in the data points would result in a significantly different model.

Bias refers to the error introduced by approximating a real-life problem, which may be extremely complicated by a much simpler model — for example, modeling non-linear problems with a linear model. In general, when we over-fit a model on given data set it results in very less bias.

This results in the variance bias trade-off.

As we fit the model over a given data set, the bias tends to decrease faster than the variance increases initially. Consequently, the expected test error (reducible) declines. However, at some point, when over-fitting starts, there is a little impact on the bias, but variance starts to increase rapidly when this happens the test error increases.

Linear Regression

Linear regression is a statistical method belonging to supervised learning used for predicting quantitative responses.

Simple Linear Regression approach predicts a quantitative response Y based on a single variable X

assuming a linear relationship. We can say :

$$\mathbb{Y} \approx \beta_0 + \beta_1 X$$

Our job is now to *estimate* β_0 and β_1 , the parameters/coefficients of our model based on the training data set, such that the hyperplane(in this case a line) is *close* to the training data set. Many criteria can estimate the closeness, the most common being *least square*.

The sum of the square of the difference between all observed response and the predicted response formulates to *Residual Sum Of Squares(RSS)*.

Problems in Linear Regression

- *Non-linearity of the response-predictor relationships.*
- *Correlation of error terms.*
- *The non-constant variance of error terms.*
- *Outliers:* when the actual prediction is very *far* from the estimated one, can arise due to inaccurate recording of data.
- *High-leverage points:* Unusual values of the predictors impact the regression line known as high leverage points.
- *Collinearity:* where two or more predictor variables are closely related to each other, it may be challenging to weed out the individual effect of a single predictor variable.

KNN Regression

KNN Regression is a non-parametric approach towards estimating or predicting values, which do not assume the form of $f(X)$. It estimates/predicts $f(x_0)$ where x_0 is a prediction point by averaging out all N_0

responses closest to x_0 . We can say:

$$f(x_0) = \frac{1}{K} \sum_{x_i \in N_0} y_i$$

Note: If K is small, the fit would be flexible and any change in the data would result in a different fit, hence for small K the variance would be high and bias low; conversely, if K is large, it might mask some structure in the data hence the bias would be high.

The Classification Problem

The responses as we discussed till now, may not always be *quantitative*, it can be also *qualitative*, predicting these qualitative responses is called classification.

We will discuss various statistical approaches to classification including:

- *SVM, Logistic Regression, KNN Classifier*
- *GAM*
- *Trees*
- *Random Forest*
- *Boosting*

Support Vector Machine(SVM)

SVM or support vector machine is the classifier that maximizes the margin. The goal of a classifier in our example below is to find a line or (n-1) dimension hyper-plane that separates the two classes present in the n-dimensional space. I have written a detailed [article](#) explaining the derivation and formulation of SVM. In my opinion, it is one of the most powerful techniques in our tool box of statistical methods in AI.

Logistic Regression

Logistic model models the probability of output response Y belonging to a particular category.

We can say:

$$\rho(X) = \frac{e^{\beta_0 + \beta_1 X}}{1 + e^{\beta_0 + \beta_1 X}}$$

Applying [componendo dividendo](#) we get:

$$\Rightarrow \frac{\rho(X)}{1 - \rho(X)} = e^{\beta_0 + \beta_1 X}$$

which is nothing but the [odds](#).

$$\Rightarrow \ln\left(\underbrace{\frac{\rho(X)}{1 - \rho(X)}}_{\log - odds / logit}\right) = \beta_0 + \beta_1 X$$

For estimating the beta coefficients, we can use [maximum likelihood](#). The basic idea is to estimate the betas such that the estimated value and observed value of the results are as close as possible. In a binary classification, with observed classes as 1 and 0 , we can say the likelihood function would look like:

$$l(\beta_0, \beta_1) = \prod_{i : y_i = 1} \hat{p}(x_i) \prod_{\tilde{i} : y_{\tilde{i}} = 0} (1 - \hat{p}(x_{\tilde{i}}))$$

Date:	17 JUNE 2020	Name:	PAVITHRAN S
Course:	R PROGRAMMING	USN:	4AL17EC068
Topic:	R PROGRAMMING	Semester & Section:	6TH B
Github Repository:	Pavithran		

AFTERNOON SESSION DETAILS

Image of session

greatlearning Learning for Life Home Live Sessions Certificates My Courses

Courses / Introduction to R / Plots_Combinations

Content

- Intro to R for Analytics Overview
- Course Overview
- Reference Material
- Introduction to R
 - R_Overview & Preliminary Steps
 - Data Types
 - Data Structures
 - Importing data
 - Data Manipulations
 - Sorting and Merging
 - Data Summaries

Plots_Combinations

```

# Income Group Distribution per Region
# Data: IncomeGroup, Region
# Frequency of Income Groups by Region

# Sample R code for the plot
library(ggplot2)
data <- read.csv("IncomeGroupDistribution.csv")
ggplot(data, aes(Region, IncomeGroup)) +
  geom_bar(position = "stack", color = "black") +
  facet_wrap(~ Region, scales = "y")

```

Report – Report can be typed or hand written for up to two pages.

Data Reshaping in R is about changing the way data is organized into rows and columns. Most of the time data processing in R is done by taking the input data as a data frame. It is easy to extract data from the rows and columns of a data frame but there are situations when we need the data frame in a format that is different from format in which we received it. R has many functions to split, merge and change the rows to columns and vice-versa in a data frame.

Joining Columns and Rows in a Data Frame

We can join multiple vectors to create a data frame using the **cbind()** function. Also we can merge two data frames using **rbind()** function.

```
# Create vector objects.
city <- c("Tampa", "Seattle", "Hartford", "Denver")
state <- c("FL", "WA", "CT", "CO")
zipcode <- c(33602, 98104, 06161, 80294)

# Combine above three vectors into one data frame.
addresses <- cbind(city, state, zipcode)

# Print a header.
cat("# # # The First data frame\n")

# Print the data frame.
print(addresses)

# Create another data frame with similar columns
new.address <- data.frame(
  city = c("Lowry", "Charlotte"),
  state = c("CO", "FL"),
  zipcode = c("80230", "33949"),
  stringsAsFactors = FALSE
)

# Print a header.
cat("# # # The Second data frame\n")

# Print the data frame.
print(new.address)

# Combine rows form both the data frames.
all.addresses <- rbind(addresses, new.address)

# Print a header.
cat("# # # The combined data frame\n")

# Print the result.
print(all.addresses)
```


When we execute the above code, it produces the following result –

```
# # # The First data frame
  city      state zipcode
[1,] "Tampa"   "FL"   "33602"
[2,] "Seattle" "WA"   "98104"
[3,] "Hartford" "CT"   "6161"
[4,] "Denver"  "CO"   "80294"

# # # The Second data frame
  city      state zipcode
1    Lowry    CO      80230
2   Charlotte FL      33949

# # # The combined data frame
  city      state zipcode
1    Tampa    FL      33602
2   Seattle  WA      98104
3   Hartford CT       6161
4    Denver  CO      80294
5    Lowry   CO      80230
6   Charlotte FL      33949
```

Merging Data Frames

We can merge two data frames by using the **merge()** function. The data frames must have same column names on which the merging happens.

In the example below, we consider the data sets about Diabetes in Pima Indian Women available in the library names "MASS". we merge the two data sets based on the values of blood pressure("bp") and body mass index("bmi"). On choosing these two columns for merging, the records where values of these two variables match in both data sets are combined together to form a single data frame.

```
library(MASS)
merged.Pima <- merge(x = Pima.te, y = Pima.tr,
  by.x = c("bp", "bmi"),
  by.y = c("bp", "bmi")
)
print(merged.Pima)
nrow(merged.Pima)
```

When we execute the above code, it produces the following result –

```
  bp  bmi npreg.x glu.x skin.x ped.x age.x type.x npreg.y glu.y
skin.y ped.y
1  60 33.8      1  117     23 0.466   27     No      2   125
20 0.088
2  64 29.7      2   75     24 0.370   33     No      2   100
23 0.368
3  64 31.2      5  189     33 0.583   29    Yes      3   158
13 0.295
```

4	64	33.2	4	117	27	0.230	24	No	1	96
27	0.289									
5	66	38.1	3	115	39	0.150	28	No	1	114
36	0.289									
6	68	38.5	2	100	25	0.324	26	No	7	129
49	0.439									
7	70	27.4	1	116	28	0.204	21	No	0	124
20	0.254									
8	70	33.1	4	91	32	0.446	22	No	9	123
44	0.374									
9	70	35.4	9	124	33	0.282	34	No	6	134
23	0.542									
10	72	25.6	1	157	21	0.123	24	No	4	99
17	0.294									
11	72	37.7	5	95	33	0.370	27	No	6	103
32	0.324									
12	74	25.9	9	134	33	0.460	81	No	8	126
38	0.162									
13	74	25.9	1	95	21	0.673	36	No	8	126
38	0.162									
14	78	27.6	5	88	30	0.258	37	No	6	125
31	0.565									
15	78	27.6	10	122	31	0.512	45	No	6	125
31	0.565									
16	78	39.4	2	112	50	0.175	24	No	4	112
40	0.236									
17	88	34.5	1	117	24	0.403	40	Yes	4	127
11	0.598									

```
age.y type.y
```

1	31	No
2	21	No
3	24	No
4	21	No
5	21	No
6	43	Yes
7	36	Yes
8	40	No
9	29	Yes
10	28	No
11	55	No
12	39	No
13	39	No
14	49	Yes
15	49	Yes
16	38	No
17	28	No

```
[1] 17
```

Melting and Casting

One of the most interesting aspects of R programming is about changing the shape of the

data in multiple steps to get a desired shape. The functions used to do this are called **melt()** and **cast()**.

We consider the dataset called ships present in the library called "MASS".

```
library(MASS)
print(ships)
```

When we execute the above code, it produces the following result –

	type	year	period	service	incidents
1	A	60	60	127	0
2	A	60	75	63	0
3	A	65	60	1095	3
4	A	65	75	1095	4
5	A	70	60	1512	6
.....					
8	A	75	75	2244	11
9	B	60	60	44882	39
10	B	60	75	17176	29
11	B	65	60	28609	58
.....					
17	C	60	60	1179	1
18	C	60	75	552	1
19	C	65	60	781	0
.....					
.....					

Melt the Data

Now we melt the data to organize it, converting all columns other than type and year into multiple rows.

```
molten.ships <- melt(ships, id = c("type", "year"))
print(molten.ships)
```

When we execute the above code, it produces the following result –

	type	year	variable	value
1	A	60	period	60
2	A	60	period	75
3	A	65	period	60
4	A	65	period	75
.....				
9	B	60	period	60
10	B	60	period	75
11	B	65	period	60
12	B	65	period	75
13	B	70	period	60
.....				

```

.....
41      A   60    service    127
42      A   60    service     63
43      A   65    service   1095
.....
.....
70      D   70    service   1208
71      D   75    service     0
72      D   75    service   2051
73      E   60    service     45
74      E   60    service     0
75      E   65    service    789
.....
.....
101     C   70    incidents     6
102     C   70    incidents     2
103     C   75    incidents     0
104     C   75    incidents     1
105     D   60    incidents     0
106     D   60    incidents     0
.....
.....

```

Cast the Molten Data

We can cast the molten data into a new form where the aggregate of each type of ship for each year is created. It is done using the **cast()** function.

```

recasted.ship <- cast(molten.ships, type+year~variable, sum)
print(recasted.ship)

```

When we execute the above code, it produces the following result –

	type	year	period	service	incidents
1	A	60	135	190	0
2	A	65	135	2190	7
3	A	70	135	4865	24
4	A	75	135	2244	11
5	B	60	135	62058	68
6	B	65	135	48979	111
7	B	70	135	20163	56
8	B	75	135	7117	18
9	C	60	135	1731	2
10	C	65	135	1457	1
11	C	70	135	2731	8
12	C	75	135	274	1
13	D	60	135	356	0
14	D	65	135	480	0
15	D	70	135	1557	13
16	D	75	135	2051	4
17	E	60	135	45	0
18	E	65	135	1226	14

